

ActPlane: Letting Agents Harness Their Own System-Level Behavior with eBPF

Paper #XXX
Anonymous Submission

Abstract

AI coding agents require harnesses to apply behavioral policies for effective and safe operation: e.g. do not leak secrets, run tests before committing. Our empirical study of 64 projects shows that while project harness instructions are declared in natural language, 64% are behavioral directives.

Yet existing agent harnesses fail to connect intent-level declarations to deterministic system-level policies, leaving a *semantic gap* and making it difficult to ensure compliance or follow rules. While 83% of directives involve system-level behavior, prompt constraints operate at the probabilistic intent level. 81% of projects require cross-event state tracking, but tool-call guardrails lose track of system-level actions. 74% of system-level rules require intent-level project or task context to evaluate, while current OS-level mechanisms expect predefined static policy and return only system events without semantic context.

To address this, we argue the policy engine should be exposed as a programmable interface, allowing agents to dynamically define and adapt their policy for their own system-level actions based on their intent context. We realize this in ActPlane, which allows agents to declare cross-event behavior as labeled information-flow control (IFC) policies under a layered policy model, observing and enforcing the agents' actions across process, file, and network boundaries with semantic feedback. We implement ActPlane with eBPF and evaluate it across empirical policies, system workloads, and coding tasks, showing improved policy compliance with low end-to-end overhead.

CCS Concepts: • **Computer systems organization** → *Embedded and cyber-physical systems*; • **Security and privacy** → *Software and application security*; • **Software and its engineering** → *Software safety*.

Keywords: AI agents, eBPF, BPF-LSM, labeled information-flow control, information-flow control, provenance, semantic feedback

ACM Reference Format:

Paper #XXX. 2026. ActPlane: Letting Agents Harness Their Own System-Level Behavior with eBPF. In *Proceedings of 2026 ACM SIGOPS Annual Technical Conference (ATC '26)*. ACM, New York, NY, USA, 15 pages.

1 Introduction

AI coding agents have become widely deployed tools for software development [13]. They operate as long-running processes with access to shells, source trees, package managers, and external APIs, autonomously writing code, running tests, and managing infrastructure across extended sessions.

As agents gain autonomy, they require harnesses to apply behavioral policies. Projects encode such policies in instruction files such as `CLAUDE.md` and `AGENTS.md` [5, 17]. Our empirical study of 64 projects extracts 2,116 statements and finds that 64% are behavioral directives, 83% of which involve system-level behavior, and 81% of projects require cross-event state tracking (§3). A directive is *per-event* if it can be checked against a single operation (“do not run `git push -force`”) and *cross-event* if it requires state across operations (“run tests before committing”—the commit is forbidden unless a test execution has been observed since the last source change). Because LLM compliance is inherently probabilistic [16], policy cannot reside inside the model [3, 26]: even cooperative agents routinely violate declared constraints through long-horizon planning errors, shell-script side effects, and opaque tool behavior [?].

Enforcing these policies faces two compounding gaps. First, a *context gap*: 74% of system-level directives are not self-contained—they reference project-specific concepts (“run the test suite”) or delegate judgment to the current task (“unless explicitly requested”) that cannot be resolved without the agent’s own project and session context (§3). A static policy authored by an external administrator cannot supply this context.

Second, a *semantic gap* between declared policy and system effects. Prompt constraints remain probabilistic, while tool-level guards [23, 27] lose coverage when agents shell out or spawn subprocesses: the same `git push` can be reached through a tool call, a `bash -c` string, or a compiled binary. OS-level policy [6, 15, 19] sees these effects but requires policies to be statically pre-written and returns opaque errors with no connection to declared intent. The agent—the entity that holds both the project context and the task intent needed to instantiate 74% of its own rules—has no programmatic interface to define or evolve the constraints governing its system-level actions.

Our insight is that the policy engine should be exposed as a programmable interface that lets agents become the control plane of their own system-level actions. Because the agent is the only entity that holds the project structure and

current task intent needed to concretize its rules, the agent must author its own enforcement policy—trusted to honestly declare intent, but not trusted to reliably follow it during execution.

ActPlane is such an interface. Agents and developers declare behavioral policies in a compact DSL that agents themselves can generate from natural-language directives; ActPlane compiles them into labeled information-flow rules and loads them into an eBPF/BPF-LSM backend. Named labels propagate across process, file, and network objects at kernel hooks, encoding execution history as checkable state: when a process reads a file labeled SECRET, the process acquires that label; when it forks, the child inherits it. When a labeled subject matches a rule, ActPlane returns semantic feedback—which rule was matched, why, and what alternative path satisfies it. Each rule independently selects notify (observe), block (pre-operation denial), or kill (post-operation termination). Policies are layered: higher-authority rules (e.g., platform-level “do not leak secrets”) cannot be weakened by the agent, while agent-authored rules can be revised as task intent evolves.

This paper makes three contributions.

1. An **empirical study** of instruction files from 64 popular projects, showing that cross-event information-flow policies are pervasive (81% of repositories) and that 74% of system-level directives require project or task context that only the agent can supply.
2. **ActPlane**, a programmable OS-level control plane for agent harnesses. ActPlane compiles behavioral policies into eBPF/BPF-LSM labeled information-flow rules, observes and enforces them across process, file, and network boundaries, returns semantic feedback that reconnects rule matches to intent-level remediation, and supports layered policies where higher-authority constraints cannot be weakened by the agent.
3. An **evaluation** of ActPlane’s policy coverage, runtime cost, and feedback effect. We measure coverage over 607 OS-enforceable directives from the empirical study, quantify per-event and end-to-end overhead, and evaluate repository-grounded policy compliance on a 20-task OctoBench subset selected for system-observable policy points against prompt-only and tool-regex baselines.

2 Background

2.1 Agent Harnesses and Behavioral Policies

An agent harness is the non-model infrastructure that mediates between an LLM agent and its execution environment [3, 26]. Harness components include tool dispatch, memory management, orchestration, observation, enforcement, and feedback. This paper focuses on the observation, enforcement, and feedback components: the mechanisms that ensure an agent’s behavioral policies hold at runtime.

Behavioral policies are constraints that govern an agent’s observable effects rather than its internal reasoning. Projects encode them in instruction files such as CLAUDE.md and AGENTS.md [5, 17]: do not push directly to main; run tests before committing; never expose secrets to the network. These policies differ from coding-style directives (e.g., “prefer const over let”) in that they produce observable effects at the system-call boundary.

2.2 The Semantic Gap

Agent policies are written as natural-language instructions plus current task intent, but an agent harness usually observes tool calls rather than the system effects those tool calls produce. This creates two mismatches. A single tool call such as `run_command("make")` can trigger hundreds of system calls. Conversely, the same forbidden effect, such as `git push`, can be reached through a direct tool call, a `bash -c` string, a Python subprocess, or a compiled helper binary. This opacity creates a *semantic gap*: policies declared in project instructions cannot be reliably enforced at only the prompt or tool-call boundary.

Prompt instructions are probabilistic: LLM compliance degrades with context length and task complexity [16]. Tool-call guards [23, 27] are deterministic but incomplete: they intercept framework actions, not the system effects that shells, scripts, and subprocesses produce. OS-level mechanisms such as syscall filters and sandboxes [6, 19] see those effects, but typically require static administrator-written policies and return opaque `-EPERM` failures with no connection to the agent’s declared rules. Existing mechanisms therefore cover different parts of the stack, but none connects instruction-level policy, tool execution, system-level effects, and semantic feedback.

A cooperative agent does not intentionally bypass tool-layer guards, but the opacity of the tool-to-effect mapping causes unintentional bypass in routine operation. Three representative scenarios illustrate this. (1) An agent calls `run_command("npm test")`; the test harness internally invokes `git push` to a staging server. The tool-call guard sees `npm test` (a permitted action), while the kernel observes the forbidden `git push` effect. (2) An agent calls `read_file(".env")` to inspect a configuration value, then later calls `run_command("curl api.example.com")`. Each tool call is individually permitted, but the combination constitutes secret exfiltration, a cross-event condition visible only to an enforcer that tracks information flow from file read to network connect. (3) An agent writes a Python script containing `subprocess.run(["git", "push", "-force"])` and executes it. The tool-call guard sees `run_command("python script.py")`; the kernel observes the nested `execve("git", ["push", "-force"])`.

3 Empirical Study

This section characterizes behavioral directives in real agent instruction files and identifies which directives describe system-level behavior, require OS-level enforcement, require cross-event state, or need project or task context.

3.1 Methodology

The study is organized around four empirical questions: how much instruction-file content is directive rather than descriptive, which topics contain those directives, which directives have system-observable counterparts, and which directives require project or task context before they can be translated into concrete rules.

Corpus construction. We collected public GitHub repositories that contain `CLAUDE.md` or `AGENTS.md` in a standard location, prioritizing projects in the AI-agent ecosystem rather than filename-only matches. We excluded non-code repositories, inactive projects, likely fake-star or SEO repositories, and trivial instruction-file stubs, but kept repositories whose instruction files were purely descriptive. The snapshot was collected on 2026-05-23 UTC and contains 64 repositories, 84 deduplicated instruction files, and 2,116 extracted statements.

Statement extraction. A statement is a coherent unit that expresses one factual claim, directive, or constraint; it need not align with a Markdown paragraph or list item. We used a two-pass LLM-assisted pipeline to extract statements with source line ranges and four labels: content type, topic, enforcement level, and context requirement. A validation script checked that annotated line ranges cover every source line exactly once and that each extracted text span matches the original file.

Coding and validation. The four labels correspond directly to the empirical questions below; each E-RQ defines the labels it uses. The second LLM pass was used for cross-validation rather than bulk keyword labeling. A stratified sample of 100 statements was independently reviewed by human annotators using the same taxonomy, with disagreements resolved by discussion. Enforceability labels were also reviewed by independent Claude and Codex agents, and the resulting disagreement reports were used to correct systematic errors such as treating tool-call directives as semantic-only or under-classifying workflow ordering constraints.

The following subsections report the resulting distributions. E-RQ3 also identifies the OS-enforceable subset used later in the evaluation.

3.2 E-RQ1: Are Instruction Files Behavioral Policies?

What we test. We first ask whether instruction files are mostly project documentation or policy. A statement is a *directive* if it asks the agent to perform, avoid, or condition an action; otherwise it is descriptive project context.




Figure 1. Directive fraction per repository by statement count. Most repositories contain a majority of directive statements.

Finding. By statement count, instruction files are primarily directive rather than descriptive. Of the 2,116 statements, 1,361 are directives (64.3%) and 755 are descriptions (35.7%). The median repository has 15 directives and a 71.0% directive fraction. This statement-level view differs from line-count analysis: directives are terse, while descriptions often contain long directory listings, command references, or architecture notes. A file can therefore look balanced by lines while still containing many independent rules the agent is expected to follow.

3.3 E-RQ2: Which Topics Contain Directives?

What we test. We next ask whether the directive/descriptive split varies by topic. We assign each statement to one of 12 topic categories adapted from prior instruction-file studies, but apply the categories at statement granularity rather than file granularity.

Finding. Directives are not evenly distributed across topics. Development Process and Implementation Details are both directive-heavy: 86.7% and 85.2% of their statements are directives, respectively. Architecture, in contrast, is mostly descriptive: directory layouts and design summaries make it one of the largest topics by lines, but only 23.0% of architecture statements are directives. This explains why file-level topic labels can be misleading for enforcement: a file classified as “Architecture” may mostly describe the project, while a shorter Development Process section may contain many enforceable rules.

3.4 E-RQ3: Which Directives Require OS-Level Enforcement?

What we test. We then ask which directives have system-observable counterparts and which require OS-level enforcement. Each directive is assigned to the first matching layer in an enforcement waterfall: *semantic-only* for reasoning,




Figure 2. Directive ratio by topic. Development Process and Implementation Details are directive-dense; Architecture is mostly descriptive.

communication, or output style; *content* for predicates requiring file-content inspection; *per-event* for predicates over a single command, file access, or network connection; and *cross-event* for predicates requiring history, ordering, lineage, or information flow. We use *system-observable* for content, per-event, and cross-event directives; we use *OS-enforceable* for the per-event and cross-event subset that ActPlane can enforce at process, file, and network hooks.

Finding. Most directives have system-observable counterparts. Of the 1,361 directives, 1,127 (82.8%) involve commands, files, network effects, or file contents. They require different enforcement layers: 234 directives (17.2%) are semantic-only, 520 (38.2%) require content inspection, 392 (28.8%) can be matched from a single OS event, and 215 (15.8%) require cross-event state. The 607 per-event and cross-event directives form the OS-enforceable subset used by ActPlane. The cumulative view shows why an OS-level IFC layer is needed: content inspection covers 38.2% of directives; adding per-event matching reaches 84.2%; the remaining 15.8% require cross-event state tracking (Figure 3). At the repository level, cross-event policies are pervasive: 81% of repositories contain at least one, and 43% require all four enforcement layers (Figure 4).

3.5 E-RQ4: What Context Is Needed to Instantiate Rules?

What we test. Finally, we ask whether system-observable directives can be instantiated as concrete rules from their text, or whether they require additional context. A directive is *self-contained* if all commands, paths, and patterns are explicit. It requires *project context* if an unresolved repository-specific concept must be resolved, such as “the full test suite” or “the migration tool.” It requires *task context* if the rule depends on the current user request or a per-session grant, such as “unless explicitly requested” or “without approval.”




Figure 3. Cumulative directive coverage by enforcement layer. Linters cover 55.4%; per-event matching reaches 84.2%; cross-event IFC closes the remaining 15.8%.




Figure 4. Fraction of repositories requiring each enforcement layer. 81% need cross-event tracking; 43% need all four layers.

Finding. Most system-observable directives are not self-contained. Of the 1,127 system-observable directives (content + per-event + cross-event), only 26.4% contain all commands, paths, and patterns needed to instantiate a concrete rule directly from the text. Another 64.2% require *project context*: the directive references a project-specific concept such as “the test suite,” “the linter,” “upstream source,” or “the migration tool,” whose concrete command or path must be resolved from the repository. The remaining 9.4% require *task context*: the directive depends on the current user request or a per-session grant, as in “unless explicitly requested,” “without approval,” or “keep changes focused” (Figure 5). This quantifies the gap between a natural-language instruction and a deployable enforcement rule: hardcoded rules can cover only the self-contained fraction, while the rest require an agent or harness component that can read the project, understand




Figure 5. Context required to instantiate system-observable directives. 73.6% require project or task context beyond the directive text.

the current task, and convert the instruction into a concrete policy before enforcement begins.

3.6 Design Requirements

The study implies four requirements for an agent harness that enforces behavioral policies at runtime.

- **OS-level observation.** Most directives involve commands, files, or network effects, so enforcement must observe behavior below the tool API.
- **Cross-event state.** A single system call is insufficient for policies such as “test before commit” or “do not leak data read from secrets.” The harness must propagate provenance across process, file, and network objects.
- **Agent-provided context.** Many directives cannot be instantiated without project structure or current task intent. The agent must be able to supply this context as a policy control plane.
- **Semantic feedback.** A bare denial tells the agent that an operation failed, but not which declared rule it matched or what compliant alternative would satisfy the rule.

ActPlane addresses these requirements with an agent-facing policy DSL, an OS-level labeled information-flow engine, and a feedback loop that connects kernel-detected rule matches back to the agent’s declared intent.

4 Design

ActPlane is a programmable OS-level control plane for agent harnesses. It compiles behavioral policies into labeled information-flow rules at OS kernel hooks, observing and enforcing agent behavior and returning semantic feedback when a rule is matched. Figure 6 summarizes the pipeline.

4.1 Threat Model and Assumptions

ActPlane targets a *cooperative but unreliable* agent: the agent is trusted to declare its task intent and policy constraints honestly, but not trusted to follow them reliably during execution. The agent may make mistakes because of long-horizon planning errors, prompt injection, shell-script side effects, or opaque tool behavior. ActPlane does not target a malicious agent that deliberately attacks the kernel, exploits the loader, or attempts to evade policy through adversarial obfuscation.

The trusted computing base consists of the eBPF programs, BPF maps, DSL compiler, runner, feedback formatter, and higher-authority policies. The agent’s process tree, including tools, shells, subprocesses, and sub-agents, is treated as the monitored domain. Policy authority and conflict resolution are handled by the layered composition model described below.

Out of scope are kernel compromise, root or CAP_BPF compromise, deliberate side channels, and semantic content checks that are not observable at process, file, or network edges. These assumptions match the harness setting we target: cooperative coding agents that need deterministic runtime guardrails for their own system-level behavior.

4.2 Design Goals

The empirical study in §3 gives four requirements: agent-provided context, OS-boundary enforcement, cross-event state, and semantic feedback. ActPlane adopts these as system goals and adds one authority goal: agent-authored rules may specialize the session policy, but they must not weaken operator or platform constraints.

4.3 Policy Lifecycle

ActPlane separates *who supplies context* from *who enforces the resulting rule*. Higher-authority policies, such as platform or operator constraints, are supplied before the agent starts and define the non-negotiable boundary of the session. Project policies come from instruction files and repository configuration. Session policies are agent-authored: before a protected task begins, the agent reads the project instructions, resolves project context such as test commands or protected paths, incorporates task context such as user approval or task scope, and writes the resulting DSL policy.

This lifecycle is the mechanism behind the “agent as control plane” claim. A directive such as “run the full test suite before committing” does not name a concrete command. A static OS policy cannot know whether the project uses `go test ./...`, `pytest`, or `pnpm test`. The agent can inspect the repository, bind the directive to the current project, and install a rule such as `block exec "git" "commit" unless after exec "go" since write "**/*.go"`. Task-scoped directives are handled similarly: a rule involving “unless explicitly requested” is generated from the current user request or approval state, not from the repository alone.




Figure 6. ActPlane pipeline. The Rust compiler lowers the policy DSL into a fixed-size binary configuration. The eBPF loader writes it into read-only data, attaches to kernel hooks, and reports rule matches as NDJSON. The runner maps rule identifiers back to human-readable feedback.

Policy adaptation is explicit. When the task changes or the agent resolves new project context, it produces a new policy document and asks the harness to compile and load it at the next session or command boundary. The currently loaded policy remains the authority until the replacement has compiled successfully. This avoids an unsafe interval in which a failed policy update silently removes enforcement.

4.4 Layered Policy Composition

ActPlane separates a pure IFC rule from the runtime binding that makes it active. A rule is a condition, an effect, and a human-readable reason; by itself, it does not define inheritance or authority. A *domain* is the runtime policy boundary for a process tree: each process maps to one domain, and each event is checked against that domain’s effective policy. Files, endpoints, and stdio are IFC objects that carry labels, not policy domains.

A binding attaches a rule to a domain as either *locked* or *default*. Locked bindings are security invariants inherited by child domains and cannot be disabled below the binding point. Default bindings are templates: a child domain may keep them, disable them for itself, or replace them with stricter local policy. Thus a child domain may add constraints, labels, gates, or narrower scope, but it cannot disable inherited locked bindings, clear inherited labels or gates, widen scope, or mutate the definition of an inherited locked rule.

ActPlane composes policies monotonically across authority layers: platform/operator rules, project-maintainer rules, and agent/session rules. Lower-authority layers may add constraints or more specific feedback, but they cannot delete, rewrite, declassify, or weaken rules introduced by a higher-authority layer. Exceptions to a higher-authority rule must

be encoded in that same layer; an agent-authored policy cannot turn a platform block into a notify, nor can it declassify labels created by the platform policy. The prototype resolves domain bindings at compile/load boundaries; the kernel receives compact compiled rule tables and does not parse YAML or DSL on the event path.

At compile time, the policy compiler merges all active layers into one rule set and preserves each rule’s authority level in metadata. The kernel evaluates the merged rule set against each event. If multiple clauses match the same event, effects are joined by severity:

kill > block > notify > allow.

Thus lower-authority policy can only make the final decision stricter. If a project rule says notify on `git push` but a platform rule says block on the same event, the operation is blocked. If two matched rules have the same effect, ActPlane reports the higher-authority rule first, then uses deterministic rule order as a tiebreaker. The feedback file may include additional matched reasons, but the enforcement decision is always the strongest joined effect.

4.5 Policy Model

ActPlane policies are object-level information-flow rules over three node types: processes (identified by pid and command name), files (keyed by path hash), and endpoints (keyed by IPv4 address and mask). The model operates at object granularity rather than byte granularity [14], trading precision for low-overhead whole-system coverage.

4.5.1 Labels and Propagation. Each node carries a 64-bit label bitmask. A *source* declaration introduces labels: an exec source labels processes that execute a matching binary, a file source labels files matching a path pattern, and an endpoint source labels network endpoints matching an address range. Let $L(x)$ denote the label bitmask of node x . Labels propagate through fixed transfer functions:

- **fork**($p \rightarrow c$): $L(c) := L(c) \cup L(p)$.
- **exec**(p, f): $L(p) := L(p) \cup L(f) \cup S(f)$, where $S(f)$ is the set of source labels matching f .
- **read**(p, f): $L(p) := L(p) \cup L(f)$.
- **write**(p, f): $L(f) := L(f) \cup L(p)$.
- **connect**(p, e): $L(e) := L(e) \cup L(p)$.

All transfer functions are monotonic: labels accumulate and are never removed by propagation. The only mechanism that removes a label is an explicit *declassify* directive in the DSL, which clears specified labels when the agent executes a designated gate binary (e.g., a redaction tool). This monotonicity bounds over-tainting: a node’s label set grows at most to the union of all source labels reachable through observed flow edges.

These transfer functions maintain a provenance invariant: if a process p carries label ℓ , then information from some


```

source AGENT = exec "claude"
source SECRET = file "**/*.env"

# Per-event: from CoplayDev/unity-mcp
rule no-push-main:
  kill exec "git" "push" "main" if AGENT
  because "Do not push to main; open a PR instead."

# Cross-event: from chenhg5/cc-connect
rule tests-before-commit:
  block exec "git" "commit"
  if AGENT unless after exec "go" since write "**/*.go"
  because "Run `go test ./...` before committing."

```

Figure 7. Two ActPlane rules drawn from real projects. The first is a per-event rule (single `execve` match); the second is a cross-event rule requiring temporal state (after. . . since).

source tagged with ℓ has reached p through observed OS-level flow edges.

4.5.2 Rules and Effects. Unlike traditional taint analysis and IFC systems that prioritize completeness (minimizing false negatives, tolerating high false positives) [7, 11], ActPlane targets *precision*: false positives that block legitimate operations disrupt the agent’s development workflow and erode trust in the harness. The design therefore favors object-granularity labels with explicit declassification gates over byte-level tracking, accepting bounded under-tainting in exchange for near-zero false positives on compliant executions.

A rule contains one or more clauses. Each clause starts with an action verb (notify, block, or kill) followed by an operation (exec, read, write, unlink, connect, or open), a target pattern, optional positional arguments, a Boolean expression over required and forbidden labels, and an optional condition (target allow-list, lineage gate, or temporal after-gate).

Each clause declares its own effect that determines how the kernel responds to a match.

Notify: the kernel emits a match event and allows the operation to proceed.

Block: a BPF-LSM hook returns `-EPERM` before the operation commits. Block is available only when the bpf LSM is active in the running kernel; it is not silently downgraded.

Kill: the kernel sends `SIGKILL` to the matching task. From a BPF-LSM hook, kill also returns `-EPERM`; from a tracepoint, kill is post-operation termination.

The composition of multiple matching rule effects is defined by the layered policy model above.

4.5.3 Policy Language Scope. The DSL is not a general-purpose mandatory access control policy. Its primitives map to recurring agent workflow patterns: “after tests,” “through

a review tool,” “data derived from untrusted input,” and “read-only sub-agent.” The grammar is small so that policy authors can translate natural-language instructions into source/sink/gate shapes without security-policy expertise.

The intermediate DSL is a deliberate design choice. The IFC engine’s backend configuration (struct `taint_config`) consists of hundreds of fixed-size rule entries, bitmask allocations, and match-kind encodings across several thousand lines of kernel code. Directly generating this backend representation from natural language (NL→backend) would require the translating agent to understand the full engine ABI—an unrealistic expectation. The DSL compresses each rule into a single declarative statement that the compiler lowers into the correct binary layout, reducing the translation task to NL→DSL (one line per rule).

4.6 System Architecture

4.6.1 Compiler and Kernel ABI. The Rust compiler parses the DSL and lowers it into a C-compatible struct `taint_config`. Lowering allocates label bits, expands Boolean expressions into required and forbidden bitmasks, lowers glob patterns into exact, prefix, suffix, or wildcard match kinds, and compiles IPv4 patterns into network/mask pairs. Human-readable rule names and reason strings (including remediation guidance) remain in Rust metadata; the kernel receives only fixed-size rule tables and emits integer rule identifiers.

This split keeps the eBPF program verifier-friendly. The kernel program iterates over fixed-size arrays in read-only data using bounded loops, copies each rule entry into a stack-local variable before matching, and uses explicit index bounds checks for argument scanning and pattern comparison.

4.6.2 Kernel State. The eBPF backend maintains five BPF hash maps:

- `ts_proc`: pid → label bitmask and lineage-gate state.
- `ts_root`, `ts_sess`: root-process and session-level temporal (after) gate state.
- `ts_file`: path hash → label bitmask.
- `ts_endp`: IPv4 address → label bitmask.

BPF map entries are keyed by tid (one entry per Linux thread) for `ts_proc` and by path hash or IPv4 address for object maps. Updates use atomic BPF helper operations; concurrent reads from different CPUs observe a consistent label bitmask per entry.

Tracepoint programs handle process lifecycle events (fork, exec, exit) and post-operation file and connect events. BPF-LSM programs attach to `bprm_check_security`, `file_open`, and `socket_connect` for pre-operation enforcement. When BPF-LSM is not active, the loader uses tracepoint mode and disables block-effect rules. In tracepoint mode, hooks fire after the operation commits; a connect tracepoint observes an already-established connection. This is why block requires BPF-LSM: only LSM hooks provide a pre-operation

```
[ActPlane] policy rule matched
rule: tests-before-commit
effect: block
event: exec git commit
reason: Run `go test ./...` before committing.
recovery: run `go test ./...`, then retry commit.
```

Figure 8. Example corrective-feedback payload appended after a kernel rule match. The hook relays this payload into the next agent turn.

verdict. Tracepoint mode supports `notify` (post-operation observation) and `kill` (post-operation task termination).

The current implementation supports up to 32 sources, 32 rules, 16 transforms, and 16 gates per policy. Rule evaluation uses `bpf_loop` for bounded iteration, which the eBPF verifier accepts for policies exceeding 100 rules.

4.6.3 Unified Event Handling. Each hook normalizes kernel arguments into a common event structure: subject pid, object kind, access kind, backend mode, target string, optional argument vector, and optional IPv4 address. A shared handler resolves the subject’s label bitmask, applies source labels for read/connect events, evaluates all supported rules, emits at most one `TAINT_VIOLATION` event for the highest-priority matching rule, and then applies propagation updates.

4.7 Corrective Feedback

The kernel match event contains the rule identifier, effect, target string, pid, and flags indicating whether the operation was blocked or the task was killed. The Rust runner maps the rule identifier to the corresponding rule name and reason string, then appends a structured payload to `.actplane/last-violation.txt`.

Agent integration uses a post-tool hook. The `actplane` feedback-hook subcommand reads new bytes from the feedback file and returns them as `additionalContext` to compatible agent runtimes [13]. The hook is a read-only relay: it does not re-evaluate policy in userspace. The kernel is the sole authority for rule matches; userspace formats the result so the agent can select an alternative execution path. A bare `-EPERM` tells the agent only that an operation failed. The feedback payload instead identifies the declared rule that fired and the compliant recovery path, which lets a co-operative agent repair the plan rather than retrying the same operation.

5 Implementation

ActPlane is implemented as a Rust driver and an eBPF program pair. The system builds a single `actplane` binary that embeds the compiled eBPF loader, extracts it at runtime, compiles the project’s policy DSL, and runs a target command under the observation and enforcement harness.

5.1 Runner and Policy Discovery

The command `actplane run - <cmd>` discovers `actplane.yaml` by searching upward from the working directory. The policy YAML contains a policy: | DSL block. Before the target process begins executing, the loader seeds the stopped target’s pid with the `AGENT` label so that all descendants inherit the agent identity. The runner prepares the feedback file (`.actplane/last-violation.txt`) and a hook state file, exports their paths to the child environment, waits for the loader to report readiness, then resumes the target. This sequencing ensures that the first target operation is observed.

5.2 BPF Backends

The eBPF program attaches to process lifecycle tracepoints (`fork`, `exec`, `exit`), syscall tracepoints for file and connect events, and BPF-LSM hooks for pre-operation enforcement. At load time, the loader checks whether the running kernel has the bpf LSM active in `/sys/kernel/security/lsm`. If not, it falls back to tracepoint mode, in which the supported effects are `notify` and `kill`; block rules are disabled because there is no pre-operation verdict path.

5.3 Limitations

Path identity. File labels are keyed by FNV-1a path hash rather than inode/device pairs. Hard links and mount binds can cause the same file to carry different labels under different paths.

Endpoint identity. Connect matching uses numeric IPv4 addresses. Hostname, DNS, SNI, and HTTP-layer matching require userspace resolution or additional instrumentation.

Argument predicates. Positional argument predicates inspect arguments after `exec` via the tracepoint argument buffer. Pre-operation blocking of argument predicates requires an argument cache before `bprm_check_security`, which is not yet implemented.

Label precision. Labels propagate at object granularity. A process that reads one byte from a labeled file acquires the full label set. This can produce over-tainting relative to byte-level systems [14]. Section 6 measures the false-positive rate and validates that declassification and gate paths prevent unnecessary harness intervention.

6 Evaluation

The evaluation uses two sources. RQ1 and RQ2 use directive-derived policies from the empirical corpus; RQ3 uses a repository-grounded OctoBench subset with case-specific policies and the official judge. The evaluation answers three questions:

RQ1 (Policy compliance) Compared with prompt-only, tool-layer, and feedback-ablation baselines, does ActPlane improve policy compliance for directive-derived policies across direct and bypass execution paths?

RQ2 (Overhead) What is the per-event and end-to-end overhead of ActPlane’s label propagation and rule checking?

RQ3 (Repository-grounded feedback) On a scaffold-aware repository-grounded coding benchmark, does ActPlane’s OS-level enforcement and semantic feedback improve official policy compliance?

6.1 Experimental Setup

Hardware and kernel. All experiments run on a single machine with an Intel Core Ultra 9 285K CPU (24 hardware cores), 125 GiB RAM, Linux 6.15.11, ext4 filesystem, and libbpf 1.x with bpf_loop support. The live enforcement runs require BPF-LSM active (bpf in /sys/kernel/security/lsm); the RQ2 overhead artifacts record the exact LSM state used for each performance run.

Agent environment. The active agent harnesses are Claude Code (via the ActPlane feedback hook) and OpenAI Codex CLI (via tool-error feedback). Exact versions are recorded in per-run metadata.

Baseline systems. All baselines use the *same* DSL rules; the difference is the enforcement layer. Tool-layer baselines are Python scripts that read the same rule.yaml and implement DSL semantics at the tool-call level. Kernel baselines are ActPlane with features selectively disabled. We compare seven systems: prompt-only, TL-1, TL-N, app-level IFC, per-event eBPF, kernel IFC without feedback, and full ActPlane. The tool-layer and app-level baselines represent AgentSpec/Progent-style guards [23, 27] and FIDES/CaMeL-style tool-call label tracking [8, 9]; the kernel baselines represent per-event eBPF enforcement and cross-event kernel IFC without semantic feedback.

6.2 Rule Set Construction

E-RQ3 identifies 607 OS-enforceable directives in the empirical corpus: 392 per-event directives and 215 cross-event directives. RQ1 uses this subset because these directives can be represented at process, file, or network hooks. Semantic-only directives have no system-observable counterpart, while content directives require file-content inspection or linting rather than OS-level IFC. The evaluation pipeline has four actors and strict separation of concerns.

Pipeline overview.

1. **Translate** (Translation LLM): reads sampled directive text, project CLAUDE.md, cloned repo, and DSL reference; produces rule.yaml. Cannot see traces or ground truth.
2. **Generate direct traces** (Trace generator—LLM or human): reads directive and project structure; produces trace_violation.jsonl and trace_compliant.jsonl, each with a ground-truth header. Cannot see rule.yaml (prevents circular bias).

3. **Generate bypass traces** (Script): programmatically wraps the triggering command from the direct trace in bash -c, python3 subprocess, and compiled-binary variants. No LLM involvement.
4. **Replay + tested LLM decision** (Eval harness): replays trace under each system, collects system response (feedback / bare error / nothing), then issues one API call to a tested LLM (weak model) for the next-action decision.
5. **Score** (Human or scorer script): compares the tested LLM’s decision against ground truth; fills correctness and outcome (TP/TN/FP/FN) in the persisted result record.

Key separations: the translation LLM and the trace generator are different actors (translation blind spots do not hide in traces); the tested LLM never sees ground truth (it acts like a real agent); bypass traces are programmatically generated (no circular evaluation risk). All systems run on the *same* agent-generated rules—translation errors are shared noise; differences between systems reflect system capability. Figure 7 shows representative per-event and cross-event rules.

6.3 RQ1: End-to-End Policy Compliance

Goal. Given a natural-language directive and a scenario (prompt + system actions), does ActPlane’s runtime enforcement and semantic feedback lead the agent to a policy-compliant final action? “Policy” means the directive-derived ActPlane rule for that scenario. This measures the runtime value of the “agent as control plane” design while keeping rule translation and trace generation as explicit artifact-generation steps.

Method. For each of the 607 OS-enforceable directives we generate exactly **two trace scenarios**: a violation trace (tool calls that should trigger the rule) and a compliant trace (correct-path tool calls that should not). Each trace is a JSONL file mirroring a real agent session: a user prompt followed by a sequence of tool calls that set up the repo state. The last tool call is omitted—the tested LLM decides it at runtime. Ground truth is determined by the combination of prompt and directive (the same tool-call sequence can be a violation or not depending on the prompt). Total: $607 \times 2 = 1,214$ trace scenarios (census over all directives).

We implement a minimal eval harness (~100 lines Python) with three phases:

1. **Deterministic trace replay under ActPlane.** All execution runs under ActPlane; labels propagate from the first tool call. For each message in the trace, tool-use blocks are executed; text blocks are preserved in context. If ActPlane fires (kill/notify), the feedback is appended to context and replay stops.

2. **Tested LLM decision step.** One API call (temperature = 0) to a weak model that reads the accumulated context (prompt + tool history + any ActPlane feedback) and writes the next normal agent response. The model never sees ground truth, the directive, or evaluator labels.
3. **Result record creation.** The harness persists an unscored result record containing the replay context, the raw LLM response, and correctness: null. Scoring is a separate pass.

Scoring methodology. A separate scoring pass (human audit or scorer script) compares the tested LLM’s final action against ground truth and fills the correctness and outcome fields:

Table 1. Outcome classification.

Action	Truth	Result
Respected directive	violation	TP
Respected directive	not violation	TN
Violated directive	violation	FN
Respected directive (wrongly)	not violation	FP

End-to-end FP/FN captures all error sources in one measurement: translation errors (wrong pattern), agent response errors (ignoring feedback), and IFC model precision (over-tainting).

Worked examples. *Example 1—TP (kill + feedback → correct recovery):* Directive: “Run tests before committing.” Trace: edit → git add → git commit → KILLED. Feedback: “Tests are stale. Re-run go test, then commit.” Tested LLM: “I’ll run go test ./. . . then commit again.” Scorer: CORRECT.

Example 2—FN (notify + agent ignores feedback): Directive: “Do not commit without user approval.” Prompt: “Review the code for bugs.” Trace: edit → git add → git commit → NOTIFY. Tested LLM: “The commit went through. Moving on.” Scorer: INCORRECT—prompt said “review,” not “commit.”

Example 3—FP (wrong translation → over-blocking): Directive: “Run tests before committing.” Prompt: “Fix and commit.” Trace: edit → go test (PASS) → git commit → KILLED (rule says pytest not go test). Tested LLM: “I already ran go test. The system wants pytest. I can’t proceed.” Scorer: INCORRECT—the agent’s test was valid; the rule was mistranslated.

Example 4—TN (compliant, no trigger): Directive: “Run tests before committing.” Prompt: “Fix, test, and commit.” Trace: edit → go test (PASS) → git commit → no trigger. Tested LLM: “Commit successful. Done.” Scorer: CORRECT.

Example 5—Kernel IFC vs. ActPlane (same rule, no feedback vs. feedback): Directive: “Run tests before committing.” Trace: edit → git commit → KILLED. Under Kernel IFC: tested LLM receives “Permission denied” → retries blindly → INCORRECT. Under ActPlane: tested LLM receives “Tests are stale.

Re-run go test.” → “I’ll run go test then commit.” → CORRECT.

Table 2. End-to-end FP/FN by enforcement level.

Level	FN	FP	Total
Per-event	TODO/392	TODO/392	392
Cross-event	TODO/215	TODO/215	215
Total	TODO/607	TODO/607	607

Table 3. End-to-end FP/FN by context requirement.

Context	FN	FP	Total
None	TODO	TODO	TODO
Project	TODO	TODO	TODO
Task	TODO	TODO	TODO

Expected results.

Methodological notes. *One decision step is the correct granularity.* The kernel operates per-syscall; each rule match is an independent decision point. A single agent decision step (try action → get feedback → decide) is the atomic unit of the feedback loop. Multi-step sessions are tested in RQ3 (OctoBench). *Why scripted setup + real agent decision.* The setup (repo state) is scripted for reproducibility and cost; the decision step runs a real agent under ActPlane to test the full feedback loop: kernel rule match → feedback → agent interprets feedback + prompt → action. *Why no gold rules for system comparison.* All systems run on the same agent-generated rules. Translation errors are shared noise; differences between systems reflect system capability.

6.4 System Comparison Across Execution Paths

ActPlane maintains policy compliance when directive-relevant behavior moves below the agent tool API or depends on cross-event state. Existing approaches fail in three common ways: tool-layer guards lose track of system-level effects (bypass), OS-level mechanisms lack cross-event state (no IFC), and kernel mechanisms return only system events without semantic context (no feedback). This sub-experiment tests all three by comparing 7 systems on the same directives.

Method. We select all rules from the **TP set** of RQ1 (rules where ActPlane produced the correct outcome on direct traces). This isolates translation noise—any difference between systems is purely architectural. For each TP rule, we take the RQ1 violation trace and programmatically generate **3 bypass variants** by replacing the triggering tool call with an indirect execution path:

bash -c The triggering command is wrapped in bash -c '<cmd>'.
 “

Python subprocess The command is wrapped in `python3 -c "import subprocess; subprocess.run([...])"`.

Compiled binary The command is embedded in a pre-built helper binary (`./commit-helper`).

Total: N_{TP} rules $\times$ 4 traces (direct + 3 bypass). Each trace uses the same JSONL format as RQ1 and the same minimal agent harness: deterministic replay $\rightarrow$ system responds $\rightarrow$ LLM handles the response (1–2 API calls per trace).

For kernel systems (ActPlane, Kernel IFC, Per-event eBPF) the trace runs under `sudo actplane run [-flags]`. For tool-layer and app-level baselines the replay executor intercepts tool calls through the Python baseline checker.

Metrics.

- **Directive compliance rate** per system, broken down by enforcement level (per-event vs. cross-event) and execution path (direct vs. bypass).
- **Bypass gap**: the difference between direct and bypass compliance per system—shows which systems lose coverage on indirect paths.
- **Feedback recovery gap**: Kernel IFC (bare `-EPERM`) vs. ActPlane (semantic feedback)—shows the value of feedback for agent recovery.

Table 4. Directive compliance rate by system and execution path.

System	Direct	Bypass	Gap
ActPlane	TODO	TODO	TODO
Kernel IFC	TODO	TODO	TODO
Per-event eBPF	TODO	TODO	TODO
App-level IFC	TODO	TODO	TODO
TL-N	TODO	TODO	TODO
TL-1	TODO	TODO	TODO

Table 5. Directive compliance rate by system and enforcement level.

System	Per-event	Cross-event	Total
ActPlane	TODO/ N_1	TODO/ N_2	TODO
Kernel IFC	TODO/ N_1	TODO/ N_2	TODO
Per-event eBPF	TODO/ N_1	TODO/ N_2	TODO
App-level IFC	TODO/ N_1	TODO/ N_2	TODO
TL-N	TODO/ N_1	TODO/ N_2	TODO
TL-1	TODO/ N_1	TODO/ N_2	TODO

Expected results. Key observation: Kernel IFC and ActPlane have identical detection, but ActPlane achieves higher compliance because semantic feedback enables agent recovery. Kernel IFC blocks but returns bare `-EPERM`; the agent retries blindly or gives up.

6.5 RQ2: Overhead

6.5.1 Macrobenchmarks (End-to-End Workloads). We replay deterministic corpus-derived agent traces and a standard Linux kernel build under no-hit ActPlane configurations. The agent trace suite contains 20 corpus-derived traces with 68 tool actions, including 20 Bash subprocesses. The Linux build runs `defconfig + vmlinux` with `make -j24` in a fresh output directory. Each workload runs three trials per configuration. Replaying fixed workloads eliminates LLM inference variance and isolates ActPlane’s runtime overhead.




Figure 9. End-to-end overhead on deterministic workloads, normalized to native execution. At 32 active no-hit rules, ActPlane adds 1.9% on agent-trace replay and 6.5% on a Linux build; at 100 rules, overhead remains below 8.4%.

6.5.2 Microbenchmarks (Per-Syscall Latency). We measure ActPlane’s per-event latency for five syscall types (fork, exec, open, write, connect) under five no-hit configurations: native, AP-1, AP-10, AP-32, and AP-100. Each configuration $\times$ syscall type runs 10K–100K iterations pinned to a single CPU core. For each trial we compute p50 and p99 latencies; Table 6 reports the median across seven trials.

Measured operations. fork: `fork() + parent waitpid()`. exec: `fork() + child execve("/bin/true") + parent waitpid()`. open: `open(path, O_RDONLY|O_CLOEXEC)` with `close()` outside the timed region. write: `write(fd, 4096)` to an already-open temp file. connect: `UDP connect(127.0.0.1:9)` on an already-created socket.

Results. The largest absolute microbenchmark cost is on open: AP-100 adds 12.8 μ s at p50 and 14.7 μ s at p99. AP-100 adds 2.6 μ s p50 overhead to connect and 0.57 μ s to write. fork and exec include scheduler and process-management overhead, so their tail latencies are noisier.

Table 6. Per-operation latency in microseconds for no-hit policies. Values are medians across seven trials.

Operation	Native p50	AP-32 p50	AP-100 p50	AP-100 p99
open	0.58	6.92	13.40	15.36
write	0.27	0.79	0.84	1.59
connect	0.58	1.98	3.17	3.90
fork	48.94	74.05	69.33	178.01
exec	248.30	314.86	317.03	490.37

Artifact note. The primary overhead policy is no-hit: it exercises steady-state label propagation and rule scanning but excludes violation reporting and userspace feedback formatting. The exact aggregate CSV/JSON artifacts are stored under `docs/rq2-performance/results/`.

6.6 RQ3: Repository-Grounded Policy Compliance (OctoBench)

Goal. RQ3 evaluates whether ActPlane improves policy compliance in long-form coding-agent tasks in real repository environments. Unlike RQ1, which isolates one directive-derived decision point, this experiment runs complete OctoBench tasks and evaluates the resulting agent trajectory with the official whole-case checklist judge.

Benchmark. OctoBench [10] is a scaffold-aware repository-grounded coding benchmark with 217 Dockerized tasks spanning multiple instruction sources: system prompt, tool schema, project files such as `CLAUDE.md` and `AGENTS.md`, memory, skills, and user queries. The benchmark reports per-check reward (passed checks divided by total checks) and a binary full-instance pass when all checks pass. We use the official whole-case evaluator unchanged: no category fallback, no external reward, and no checklist override.

Subset and conditions. ActPlane cannot observe purely semantic checks (e.g., tone or summary quality), so we evaluate a 20-task OctoBench subset selected for system-observable policy points, with case-specific policies. Each task runs under three paired conditions: baseline (the original task with no enforcement), tool-regex (a case-specific tool-call hook), and ActPlane (case-specific OS hooks plus feedback). No shared guardrail file is used. Tool-regex intercepts only tool calls, while ActPlane observes OS-level exec, open, write, and connect edges inside the task container and injects corrective feedback when an OS-level rule fires.

Metrics. We treat the official OctoBench reward as the benchmark’s policy-compliance metric: it measures the fraction of checklist policies judged satisfied for a whole task. We also report three diagnostic submetrics computed from the same official checklist results: *User-query reward* restricts checks to the user-task category; *implementation/test reward* restricts checks whose type is implementation, testing, configuration, or modification; and *compliance reward* restricts

checks typed as compliance. These submetrics are still policy-compliance metrics backed by the official checklist judge, not deterministic test-script completion.




Figure 10. OctoBench 20-task selected subset. ActPlane improves official policy compliance over both baseline and tool-regex.

Result. On the selected 20-task subset, official reward is 0.788 for baseline, 0.818 for tool-regex, and 0.888 for ActPlane. ActPlane therefore raises official reward by 10.0 percentage points over baseline and 7.0 points over tool-regex. The improvement is not limited to generic compliance: user-query reward increases by 28.2 points over baseline, and implementation/test reward increases by 25.0 points. This supports the narrower claim that OS-level feedback can improve policy compliance in repository-grounded coding agents. We do not claim deterministic task completion from this benchmark because OctoBench is judged by an LLM checklist evaluator rather than by per-task test scripts.

7 Related Work

7.1 In-Kernel Provenance and Information Flow

PASS introduced provenance as a storage-layer primitive [18]. Hi-Fi and LPM used LSM hooks to capture trustworthy whole-system provenance [2, 21]. CamFlow extended this to a streaming provenance framework over processes, files, and sockets [20].

CamQuery is the closest prior system [19]. It executes provenance queries inline in the kernel, propagates labels across process, file, and network edges, and can deny operations that match a policy rule. ActPlane shares this label-propagation model but differs in three respects: it targets cooperative AI agents rather than adversarial intrusions, compiles an agent-oriented source/sink DSL rather than general provenance queries, and returns semantic feedback to the matching actor. The implementation differs as well: ActPlane uses the eBPF/BPF-LSM substrate rather than a loadable kernel module.

TaintDroid, Dytan, libdft, and Panorama established dynamic taint tracking at runtime, binary-instrumentation, and emulation layers [7, 11, 14, 28]. ActPlane operates at object

granularity rather than byte granularity, trading precision for low-overhead whole-system coverage.

7.2 eBPF and LSM Enforcement

BPF-LSM, originally proposed as KRSI, allows eBPF programs to attach to Linux security hooks and return allow/deny verdicts [24, 25]. ActPlane uses BPF-LSM as its pre-operation enforcement backend. Contemporary eBPF enforcement tools such as Tetragon, Tracee, and eBPF-PATROL provide per-event predicate matching and, in Tetragon’s case, a boolean lineage flag (`followChildren`) that propagates across fork/exec. These systems evaluate per-event predicates or single-channel lineage flags rather than cross-channel labeled information flow [1, 6, 12].

7.3 Agent Guardrails and Information-Flow Control

AgentSpec and Progent provide deterministic runtime enforcement over tool-call names, arguments, and framework actions [23, 27]. These systems enforce at the tool-call boundary; ActPlane complements them by enforcing below the tool API, where effects that bypass the framework are still visible.

FIDES and CaMeL apply typed information-flow control and capability separation within the agent loop or scaffolding [8, 9]. ActPlane applies the same style of information-flow reasoning at the OS boundary, where observation and enforcement persist across context windows and cannot be circumvented by subprocess indirection.

7.4 Agent Instruction Files

Recent empirical studies characterize CLAUDE.md, AGENTS.md, and related agent instruction files along dimensions of content taxonomy, maintenance practices, and effect on task efficiency [4, 5, 17, 22]. These studies classify at file or section-heading granularity. ActPlane’s empirical study (§3) classifies at statement level along four axes (content type, topic, enforcement level, context requirement), asking which instructions constitute cross-event information-flow policies, which require project or task context to instantiate, and which can be enforced at the OS level.

8 Conclusion

Our empirical study of 64 projects reveals that agent instruction files are behavioral policies (64% directives), that most directives describe system-level behavior (83% system-observable), and that 74% of system-observable directives cannot be instantiated without project or task context that only the agent possesses. These findings motivate a design in which the agent authors its own enforcement policy.

ActPlane provides a programmable OS-level control plane for agent harnesses. It compiles agent-declared behavioral policies into eBPF/BPF-LSM labeled information-flow rules, propagates labels across process, file, and endpoint objects

at kernel hooks, and returns kernel-detected rule matches as semantic feedback to the agent. A layered policy model ensures that higher-authority constraints cannot be weakened by the agent, while agent-authored rules can evolve with task intent. By supporting notify, pre-operation denial, and harness-level termination with structured remediation, ActPlane enables agents to recover after policy rule matches rather than fail opaquely. ActPlane bridges the context gap and the semantic gap that single-level harnesses leave open, unifying labeled information-flow observation and enforcement with an agent-oriented policy language and a feedback loop on the eBPF/BPF-LSM substrate.

References

- [1] Aqua Security. 2026. Tracee: Linux Runtime Security and Forensics using eBPF. <https://github.com/aquasecurity/tracee>.
- [2] Adam Bates, Dave Tian, Kevin R. B. Butler, and Thomas Moyer. 2015. Trustworthy Whole-System Provenance for the Linux Kernel. In *24th USENIX Security Symposium (USENIX Security 15)*. USENIX Association, Washington, DC, 319–334. <https://www.usenix.org/conference/usenixsecurity15/technical-sessions/presentation/bates>
- [3] Birgitta Boeckeler. 2026. Harness Engineering for Coding Agent Users. <https://martinfowler.com/articles/harness-engineering.html>.
- [4] Worawalan Chatlatanagulchai, Hao Li, Yutaro Kashiwa, Brittany Reid, Kundjanasith Thonglek, Pattara Leelaprute, Arnon Rungsawang, Bundit Manaskasemsak, Bram Adams, Ahmed E. Hassan, and Hajimu Iida. 2025. Agent READMEs: An Empirical Study of Context Files for Agentic Coding. arXiv:2511.12884. <https://arxiv.org/abs/2511.12884>
- [5] Worawalan Chatlatanagulchai, Kundjanasith Thonglek, Brittany Reid, Yutaro Kashiwa, Pattara Leelaprute, Arnon Rungsawang, Bundit Manaskasemsak, and Hajimu Iida. 2025. On the Use of Agentic Coding Manifests: An Empirical Study of Claude Code. arXiv:2509.14744. <https://arxiv.org/abs/2509.14744>
- [6] Cilium Project. 2026. Tetragon: eBPF-based Security Observability and Runtime Enforcement. <https://tetragon.io/>.
- [7] James Clause, Wanchun Li, and Alessandro Orso. 2007. Dytan: A Generic Dynamic Taint Analysis Framework. In *Proceedings of the 2007 International Symposium on Software Testing and Analysis*. Association for Computing Machinery, London, United Kingdom, 196–206. doi:10.1145/1273463.1273490
- [8] Manuel Costa, Boris Koepf, Aashish Kolluri, Andrew Paverd, Mark Russinovich, Ahmed Salem, Shruti Tople, Lukas Wutschitz, and Santiago Zanella-Beguelin. 2025. Securing AI Agents with Information-Flow Control. arXiv:2505.23643. <https://arxiv.org/abs/2505.23643>
- [9] Edoardo Debenedetti, Ilia Shumailov, Tianqi Fan, Jamie Hayes, Nicholas Carlini, Daniel Fabian, Christoph Kern, Chongyang Shi, Andreas Terzis, and Florian Tramer. 2025. Defeating Prompt Injections by Design. arXiv:2503.18813. <https://arxiv.org/abs/2503.18813>
- [10] Ding et al. 2026. OctoBench: Benchmarking Scaffold-Aware Instruction Following in Repository-Grounded Agentic Coding. arXiv:2601.10343.
- [11] William Enck, Peter Gilbert, Byung-Gon Chun, Landon P. Cox, Jaeyeon Jung, Patrick McDaniel, and Anmol N. Sheth. 2010. TaintDroid: An Information-Flow Tracking System for Realtime Privacy Monitoring on Smartphones. In *9th USENIX Symposium on Operating Systems Design and Implementation (OSDI 10)*. USENIX Association, Vancouver, Canada, 393–407. <https://www.usenix.org/conference/osdi10/taintdroid-information-flow-tracking-system-realtime-privacy-monitoring>
- [12] Sangam Ghimire, Nirjal Bhurtel, Roshan Sahani, and Sudan Jha. 2025. eBPF-PATROL: Protective Agent for Threat Recognition and Overreach Limitation using eBPF in Containerized and Virtualized Environments. arXiv:2511.18155. <https://arxiv.org/abs/2511.18155>
- [13] Yuxuan Jiang, Meng Xu, Yiwen Song, and Xinwen Fu. 2025. AgentSight: Bridging the Semantic Gap Between Agent Intent and Low-Level Actions. arXiv:2508.02736. <https://arxiv.org/abs/2508.02736>
- [14] Vasileios P. Kemerlis, Georgios Portokalidis, Kangkook Jee, and Angelos D. Keromytis. 2012. libdft: Practical Dynamic Data Flow Tracking for Commodity Systems. In *Proceedings of the 8th ACM SIGPLAN/SIGOPS Conference on Virtual Execution Environments*. Association for Computing Machinery, London, United Kingdom, 121–132. doi:10.1145/2151024.2151042
- [15] Maxwell Krohn, Alexander Yip, Micah Brodsky, Natan Cliffer, M. Frans Kaashoek, Eddie Kohler, and Robert Morris. 2007. Information Flow Control for Standard OS Abstractions. In *Proceedings of the 21st ACM SIGOPS Symposium on Operating Systems Principles (SOSP '07)*. ACM, 321–334.
- [16] Nelson F. Liu, Kevin Lin, John Hewitt, Ashwin Paranjape, Michele Bevilacqua, Fabio Petroni, and Percy Liang. 2024. Lost in the Middle: How Language Models Use Long Contexts. In *Transactions of the Association for Computational Linguistics*, Vol. 12. 157–173.
- [17] Jai Lal Lulla, Seyedmoein Mohsenimofidi, Matthias Galster, Jie M. Zhang, Sebastian Baltes, and Christoph Treude. 2026. On the Impact of AGENTS.md Files on the Efficiency of AI Coding Agents. arXiv:2601.20404. <https://arxiv.org/abs/2601.20404>
- [18] Kiran-Kumar Muniswamy-Reddy, David A. Holland, Uri Braun, and Margo Seltzer. 2006. Provenance-Aware Storage Systems. In *2006 USENIX Annual Technical Conference (USENIX ATC 06)*. USENIX Association, Boston, MA, 43–56. <https://www.usenix.org/conference/2006-usenix-annual-technical-conference/provenance-aware-storage-systems>
- [19] Thomas F. J.-M. Pasquier, Xueyuan Han, Thomas Moyer, Adam Bates, Olivier Hermant, David Eysers, Jean Bacon, and Margo Seltzer. 2018. Runtime Analysis of Whole-System Provenance. In *Proceedings of the 2018 ACM SIGSAC Conference on Computer and Communications Security*. Association for Computing Machinery, Toronto, Canada, 1601–1616. doi:10.1145/3243734.3243776
- [20] Thomas F. J.-M. Pasquier, Jatinder Singh, David Eysers, and Jean Bacon. 2017. CamFlow: Managed Data-Sharing for Cloud Services. *IEEE Transactions on Cloud Computing* 5, 3 (2017), 472–484. doi:10.1109/TCC.2015.2489211
- [21] Devin J. Pohly, Stephen McLaughlin, Patrick McDaniel, and Kevin Butler. 2012. Hi-Fi: Collecting High-Fidelity Whole-System Provenance. In *Proceedings of the 28th Annual Computer Security Applications Conference*. Association for Computing Machinery, Orlando, FL, 259–268. doi:10.1145/2420950.2420989
- [22] Helio Victor F. Santos, Vitor Costa, Joao Eduardo Montandon, and Marco Tulio Valente. 2025. Decoding the Configuration of AI Coding Agents: Insights from Claude Code Projects. arXiv:2511.09268. <https://arxiv.org/abs/2511.09268>
- [23] Tianneng Shi, Jingxuan He, Zhun Wang, Hongwei Li, Linyu Wu, Wenbo Guo, and Dawn Song. 2025. Progent: Programmable Privilege Control for LLM Agents. arXiv:2504.11703. <https://arxiv.org/abs/2504.11703>
- [24] KP Singh. 2019. Kernel Runtime Security Instrumentation. Linux Plumbers Conference. <https://lpc.events/event/4/contributions/454/>
- [25] The Linux Kernel Documentation. 2021. LSM BPF Programs. https://www.kernel.org/doc/html/v5.15/bpf/bpf_lsm.html
- [26] Vivek Trivedy. 2026. The Anatomy of an Agent Harness. <https://www.langchain.com/blog/the-anatomy-of-an-agent-harness>
- [27] Haoyu Wang, Christopher M. Poskitt, and Jun Sun. 2025. AgentSpec: Customizable Runtime Enforcement for Safe and Reliable LLM Agents. arXiv:2503.18666. <https://arxiv.org/abs/2503.18666>

- [28] Heng Yin, Dawn Song, Manuel Egele, Christopher Kruegel, and Engin Kirda. 2007. Panorama: Capturing System-wide Information Flow for Malware Detection and Analysis. In *Proceedings of the 14th ACM*

Conference on Computer and Communications Security. Association for Computing Machinery, Alexandria, VA, 116–127. doi:10.1145/1315245.1315261